

BigData Processing with Apache Spark

Alfred Thompson

Jul 20, 2016



Approach

- ▶ Part 1: Overview
- ▶ Part 2: Demo

Overview

PART ONE

What is BigData?

- ▶ huge datasets
- ▶ overwhelms single node, traditional clusters
- ▶ multiple dimensions of analysis
- ▶ ever growing
- ▶ critical use cases
- ▶ go cheap to scale out big

Data as Trash

- it scales...



Figure: litter



Figure: trash heaps



Figure: garbage cities



Figure: oceanic gyres

Data as Exhaust

- ▶ logs thrown off from system events, where system is
 - ▶ server
 - ▶ edge device
 - ▶ router/switch
 - ▶ mobile phone/watch/tablet
 - ▶ media players and servers
 - ▶ RFID transactions
- ▶ stuff to get
 - ▶ search terms
 - ▶ indexing logs
 - ▶ error logs
- ▶ can be correlated
 - ▶ *term* was searched from *geo* running *OS*
 - ▶ *phone* belonging to *user* entered *store*
 - ▶ *pasta and soda* were bought at *timestamp* from *vendor*

Conspiracy of Bigness



Figure: microprocessor

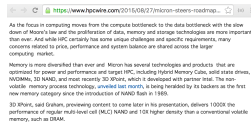


Figure: memory

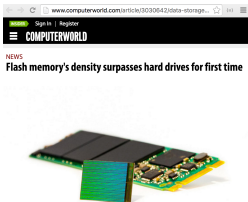


Figure: storage

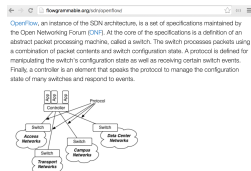


Figure: network

What is Grid/Cluster?

- ▶ gangs of computers
- ▶ HA: high availability
 - ▶ focus on nonstop processing
 - ▶ use fewer, higher end hardware nodes
 - ▶ multipath networking and I/O channels
- ▶ HPC: high performance computing
 - ▶ focus on max throughput running massive workloads
 - ▶ run many, dense capacity nodes
 - ▶ fortran is the killer app
 - ▶ message passing
- ▶ deployment models
 - ▶ shared memory
 - ▶ shared nothing
 - ▶ master-slave
 - ▶ multi-master

Grids as Trash

- ▶ build grids with off-the-shelf hardware
 - ▶ commodity servers
 - ▶ low cost
 - ▶ bulk availability
 - ▶ assume failure, be resilient when disaster comes
 - ▶ virtualizable
 - ▶ replication
- ▶ forerunners: beowulf, KLAT2

Other Conspiracies

- ▶ functional programming
- ▶ JVM/Groovy/Scala/Clojure
- ▶ open source
- ▶ mobile platforms
- ▶ social networks
- ▶ e-commerce
- ▶ advertising
- ▶ data science

Seminal Efforts

- ▶ Google Papers
 - ▶ GFS
 - ▶ MapReduce
- ▶ Yahoo!/Nutch/Hadoop

Hadoop ecosystem

- ▶ HDFS
- ▶ MapReduce
- ▶ YARN
- ▶ HBase
- ▶ Zookeeper
- ▶ plus OLAP, SQLs, graph, event processing, job workflow

Next Stage of Evolution: Spark

- ▶ high speed, in memory processing
- ▶ simplified programming model
- ▶ toolkit for machine learning
- ▶ efficiently stream/spill to I/O
- ▶ reuse, extend the Hadoop ecosystem

Spark Architecture

- ▶ RDD: resilient distributed dataset
- ▶ partition
- ▶ accumulate transformations over dataset
- ▶ process lazily
- ▶ program in Scala, Python, Java
- ▶ process data as SQL tables, graphs, streams
- ▶ deploys standalone, hadoop, mesosphere, AWS

Use Cases

- ▶ batch: better mapreduce
- ▶ machine learning
- ▶ data science
- ▶ complex event processing

Break

INTERMISSION

PART TWO

Flow

- ▶ create a VPC with internet gateway & subnet with external egress
- ▶ create a security group allowing SSH ingress
- ▶ start an EMR cluster with Spark and Zeppelin installed
- ▶ SSH into cluster master to verify install
- ▶ SSH tunnel notebook port to localhost
- ▶ navigate to notebook port and run Spark operations
- ▶ save analysis (as needed)
- ▶ terminate cluster

Flight Data

- ▶ access zeppelin
- ▶ create note header
- ▶ load data
- ▶ sample data and load to table
- ▶ query: 10 airports with most departures
- ▶ query: most flight delays over 15 mins
- ▶ query: most flight delays over 60 mins
- ▶ query: most flight cancellations
- ▶ query: most popular flight routes

Access Zeppelin

- ▶ navigate to <http://localhost:8890/>

Create Note Header

```
%md
Flight delays and cancellations for domestic flights in the U.S.
===
[Al Thompson] (http://github.com/almacro)
---
```

Reference:

```
[New - Apache Spark on Amazon EMR]
(https://aws.amazon.com/blogs/aws/new-apache-spark-on-amazon-emr)
```

Load Data

```
val parquetFile = sqlContext
  .read
  .parquet("s3://us-east-1.elasticmapreduce.samples/flightdata/input")
parquetFile.count()  /** Number of rows */
```

Sample and Load Data

```
parquetFile.take(3).foreach(println)  
parquetFile.registerTempTable("flights")
```

Query: Most Departures

%md

The 10 airports with the most departures

===

%sql

SELECT origin, count(*) AS total_departures FROM flights

WHERE year >= 2000

GROUP BY origin

ORDER BY total_departures

DESC LIMIT 10

Query: Most 15+ min Delays

%md

The most flight delays over 15 minutes

===

%sql

SELECT origin, count(depdelayminutes) AS total_delays

FROM flights

WHERE depdelayminutes > 15 and year >=2000

GROUP BY origin

ORDER BY total_delays

DESC LIMIT 10

Query: Most 60+ min Delays

%md

The most flight delays over 60 minutes

===

%sql

SELECT origin, count(depdelayminutes) AS total_delays

FROM flights

WHERE depdelayminutes > 60 and year >=2000

GROUP BY origin

ORDER BY total_delays

DESC LIMIT 10

Query: Most Cancellations

%md

The most flight cancellations

===

%sql

SELECT origin, count(cancelled) AS total_cancellations

FROM flights

WHERE cancelled = 1 and year >=2000

GROUP BY origin

ORDER BY total_cancellations

DESC LIMIT 10

Query: Most Cancellations of Total Flights

%md

The most flight cancellations out of total flights, as percentage

===

%sql

```
with total_flights as (  
  SELECT origin, count(*) AS total_flights  
  FROM flights  
  WHERE year >=2000  
  GROUP BY origin  
  ORDER BY total_flights  
  DESC LIMIT 10), total_cancellations as  
(  
  SELECT origin, count(cancelled) AS total_cancellations  
  FROM flights  
  WHERE cancelled = 1 and year >=2000  
  GROUP BY origin  
  ORDER BY total_cancellations  
  DESC LIMIT 10)  
SELECT flights.origin,  
       total_flights,  
       total_cancellations,  
       format_number(100*total_cancellations/total_flights, 2)  
       AS percent  
FROM total_flights flights  
INNER JOIN total_cancellations cancellations  
ON flights.origin = cancellations.origin  
ORDER BY percent DESC
```

Query: Most Popular Routes

%md

The 10 most popular flight routes

===

%sql

SELECT origin, dest, count(*) AS total_flights FROM flights

WHERE year >= 2000

GROUP BY origin, dest

ORDER BY total_flights

DESC LIMIT 10

Resources

- ▶ aws blog

<https://aws.amazon.com/blogs/aws/new-apache-spark-on-amazon-emr/>

- ▶ cloud academy

<http://cloudacademy.com/blog/big-data-amazon-emr-apache-spark-and-apache-zeppelin-part-one-of-two/>

<http://cloudacademy.com/blog/big-data-getting-started-with-amazon-emr-apache-spark-and-apache-zeppelin-part-two-of-two/>

- ▶ spark project

<http://spark.apache.org/>

- ▶ zeppelin project

<https://zeppelin.apache.org/>

Thanks

PROCESS COMPLETED

Extras

EXTRA INNINGS

Create Gateway

```
aws ec2 create-internet-gateway
```

Create VPC

```
aws ec2 create-vpc --cidr-block 10.101.0.0/16
```

Configure DNS in VPC

```
aws ec2 modify-vpc-attribute \  
  --vpc-id AAA \  
  --enable-dns-support "{\"Value\": true}"
```

```
aws ec2 modify-vpc-attribute \  
  --vpc-id AAA \  
  --enable-dns-hostnames "{\"Value\": true}"
```

Configure Networking

```
aws ec2 attach-internet-gateway \  
  --internet-gateway-id AAA \  
  --vpc-id BBB  
  
aws ec2 create-subnet \  
  --vpc-id AAA \  
  --cidr-block 10.101.1.0/24  
  
aws ec2 create-security-group \  
  --name sparkdemo-sg \  
  --description "Spark Demo, allow inbound" \  
  --vpc-id AAA  
  
aws ec2 authorize-security-group-ingress \  
  --group-id AAA \  
  --protocol tcp \  
  --port 22 \  
  --cidr 0.0.0.0/0
```

Configure Routes

```
aws ec2 create-route-table --vpc-id AAA
```

```
aws ec2 create-route \  
  --route-table-id AAA \  
  --gateway-id BBB \  
  --destination-cidr-block 0.0.0.0/0
```

```
aws ec2 associate-route-table \  
  --subnet-id AAA \  
  --route-table-id BBB
```

Create Cluster

```
aws emr create-cluster \  
  --name "Demo Spark Cluster" \  
  --release-label emr-4.7.2 \  
  --applications \  
    Name=Spark Name=Zeppelin-Sandbox \  
  --ec2-attributes \  
    KeyName=AAA,SubnetId=BBB,AdditionalMasterSecurityGroups=CCC \  
  --instance-type m4.xlarge \  
  --instance-count 3 \  
  --use-default-roles
```

Describe Cluster

```
aws emr describe-cluster --cluster-id AAA
```

SSH to Master

```
aws emr ssh --cluster-id AAA --key-pair-file BBB
```


Verify Install

`spark-shell`

`pyspark`

`sparkR`

SSH Tunnel for Notebook port

```
ssh \
-o StrictHostKeyChecking=no \
-o ServerAliveInterval=10 \
-i AAA \
-L 8890:127.0.0.1:8890 \
hadoop@ec2-XXX.compute-1.amazonaws.com \
-Nv
```

Terminate Cluster

```
aws emr terminate-clusters --cluster-ids AAA
```

All Done



Logout